


线性回归 Linear Regression

- Given the training dataset of (data,label) pairs,

$$D = \{(x^{(i)}, y^{(i)})\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label

$$y^{(i)} \approx f_{\theta}(x^{(i)})$$


$$f_{\theta}(x) = \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n + \theta_0$$

- Function set $\{f_{\theta}(x^{(i)})\}$ is called hypothesis space
- Learning is referred to as updating the parameter θ to make the prediction closed to the corresponding label

逻辑斯蒂回归

Logistic Regression

- Given the training dataset of (data,label) pairs,

$$D = \{(x^{(i)}, y^{(i)})\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label

$$y^{(i)} \approx f_{\theta}(x^{(i)}) \quad \rightarrow \quad f_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

- Function set $\{f_{\theta}(x^{(i)})\}$ is called hypothesis space
- Learning is referred to as updating the parameter θ to make the prediction closed to the corresponding label

决策树学习

Decision Tree Learning


- Given the training dataset of (data,label) pairs,

$$D = \{(x^{(i)}, y^{(i)})\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label

$$y^{(i)} \approx f_{\theta}(x^{(i)}) \Rightarrow$$


a decision tree



- Function set $\{f_{\theta}(x^{(i)})\}$ is called **hypotheses**
- Learning is referred to as updating the parameter θ to make the prediction closed to the corresponding label

例子 Computer Vision: Car Detection

You see this:

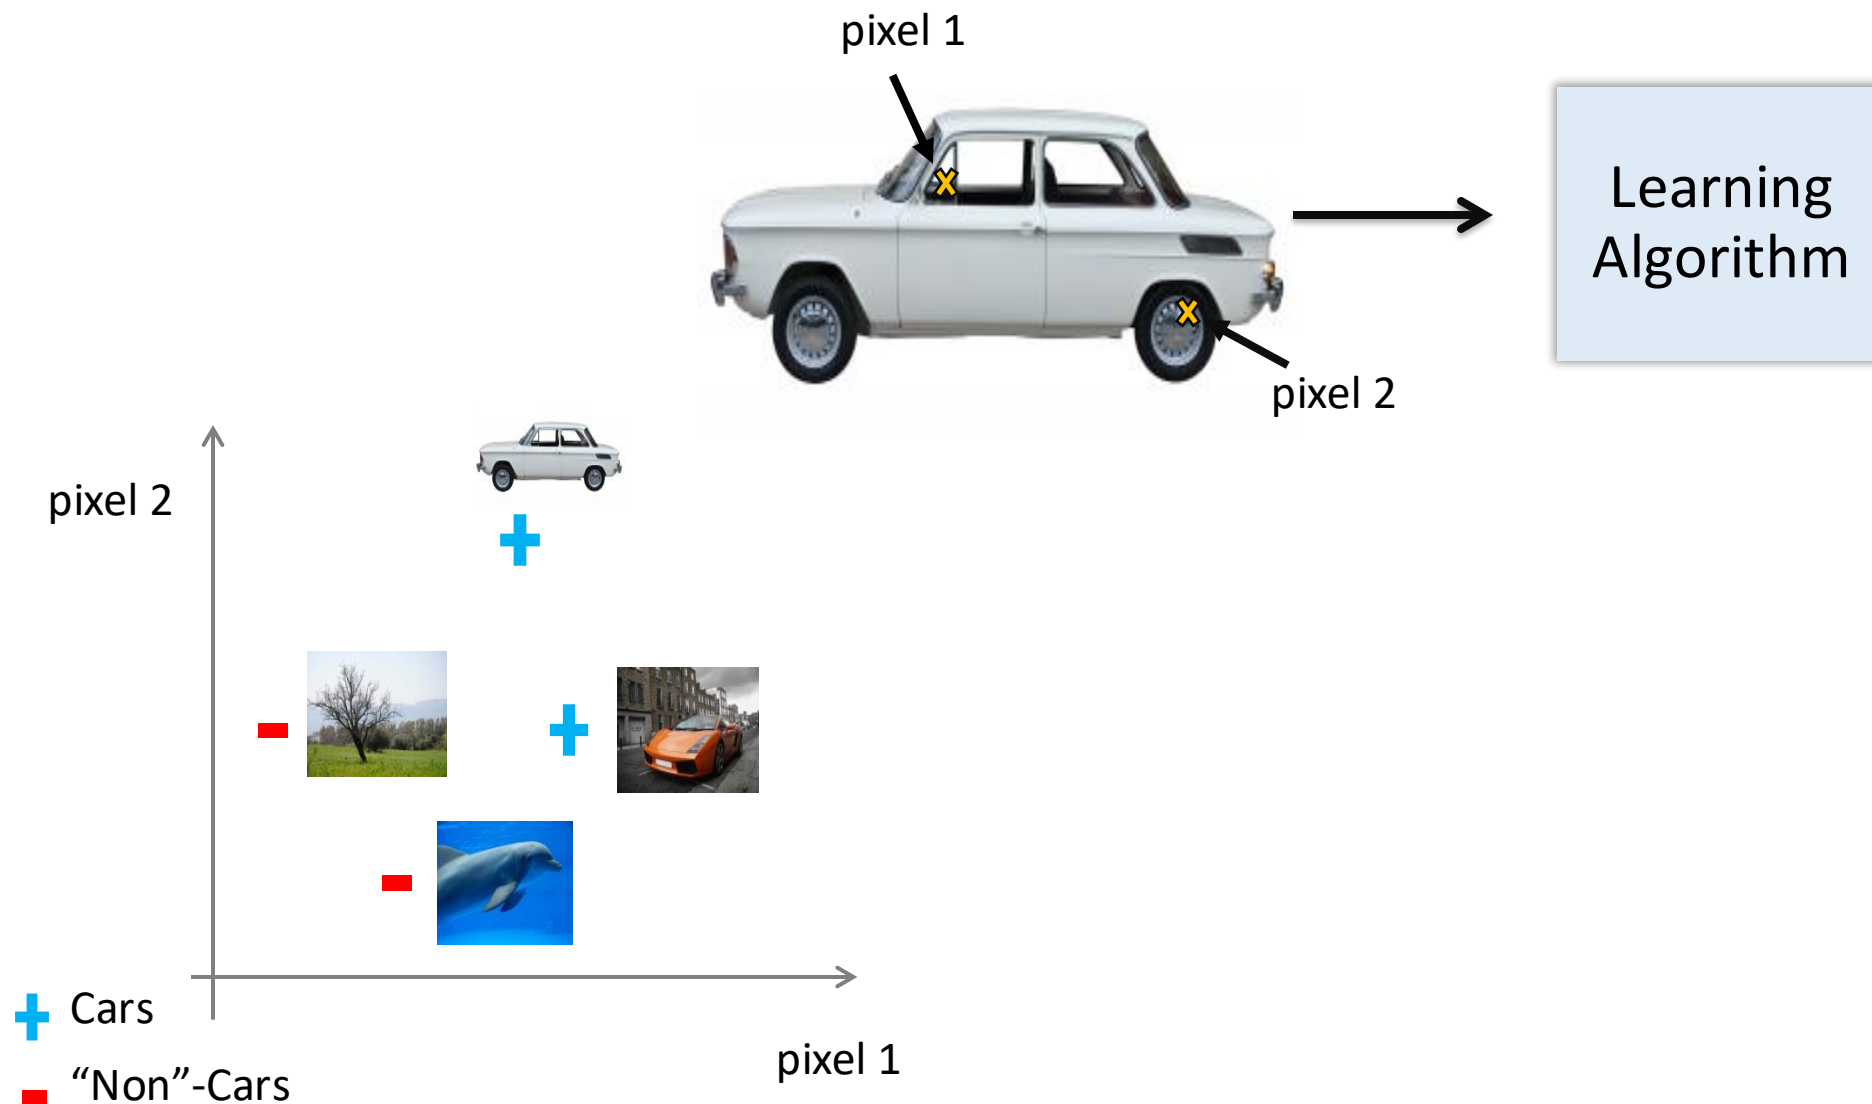


What is this?


But the camera sees this:

194	210	201	212	199	213	215	195	178	158	182	209
180	189	190	221	209	205	191	167	147	115	129	163
114	126	140	188	176	165	152	140	170	106	78	88
87	103	115	154	143	142	149	153	173	101	57	57
102	112	106	131	122	138	152	147	128	84	58	66
94	95	79	104	105	124	129	113	107	87	69	67
68	71	69	98	89	92	98	95	89	88	76	67
41	56	68	99	63	45	60	82	58	76	75	65
20	43	69	75	56	41	51	73	55	70	63	44
50	50	57	69	75	75	73	74	53	68	59	37
72	59	53	66	84	92	84	74	57	72	63	42
67	61	58	65	75	78	76	73	59	75	69	50

例子 Computer Vision: Car Detection




例子 Computer Vision: Car Detection




大脑的学习算法 Brain Learning Algorithm

The “one learning algorithm” hypothesis



Somatosensory cortex learns to see

大脑中的神经元 Neurons in the Brain





- The brain is composed of a **mass of interconnected neurons** (about 10^{11} neurons with an average of 10^4 connections each).
 - each neuron is connected to many other neurons

- Neural Networks can be :
 - **Biological** models
 - **Artificial** models
- Desire to produce **artificial systems** capable of sophisticated computations **similar** to the human brain.

神经网络的历史

Brief History of Neural Networks

- The First wave
 - 1943 McCulloch and Pitts proposed the **McCulloch-Pitts neuron model**.
 - 1958 Rosenblatt introduced the simple single layer networks now called **Perceptrons**.
 - 1969 Minsky and Papert's book Perceptrons demonstrated the limitation of single layer perceptrons, and almost the whole field went into hibernation.




Rosenblatt and perceptron

神经网络的历史

Brief History of Neural Networks

- The Second wave
 - 1982 Energy Based Model (EBM), the hopfield network
 - 1984 the traveling salesman problem (abbr. TSP), one of the NP-complete combinatorial optimization problems, can be solved by Hopfield neural network



J. J. Hopfield, "Neural networks and physical systems with emergent collective computational abilities",
Proceedings of the National Academy of Sciences of the USA, vol. 79 no. 8 pp. 2554–2558, April 1982.

神经网络的历史

Brief History of Neural Networks

- The Second wave
 - 1986 The **Back-Propagation learning algorithm** for Multi-Layer Perceptrons was rediscovered and the whole field took off again.



David Rumelhart



Geoffery Hinton

Learning representations by back-propagating errors

David E. Rumelhart*, Geoffrey E. Hinton†
& Ronald J. Williams*

* Institute for Cognitive Science, C-015, University of California,
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,
Pittsburgh, Philadelphia 15213, USA

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure¹.

There have been many attempts to design self-organizing neural networks. The aim is to find a powerful synaptic modification rule that will allow an arbitrarily connected neural network to develop an internal structure that is appropriate for a particular task domain. The task is specified by giving the desired state vector of the output units for each state vector of

more difficult when we introduce hidden units whose actual or desired states are not specified by the task. (In perceptrons, there are 'feature analysers' between the input and output that are not true hidden units because their input connections are fixed by hand, so their states are completely determined by the input vector: they do not learn representations.) The learning procedure must decide under what circumstances the hidden units should be active in order to help achieve the desired input-output behaviour. This amounts to deciding what these units should represent. We demonstrate that a general purpose and relatively simple procedure is powerful enough to construct appropriate internal representations.

The simplest form of the learning procedure is for layered networks which have a layer of input units at the bottom; any number of intermediate layers; and a layer of output units at the top. Connections within a layer or from higher to lower layers are forbidden, but connections can skip intermediate layers. An input vector is presented to the network by setting the states of the input units. Then the states of the units in each layer are determined by applying equations (1) and (2) to the connections coming from lower layers. All units within a layer have their states set in parallel, but different layers have their states set sequentially, starting at the bottom and working upwards until the states of the output units are determined.

The total input, x_j , to unit j is a linear function of the outputs, y_i , of the units that are connected to j and of the weights, w_{ji} , on these connections


$$x_j = \sum_i y_i w_{ji} \quad (1)$$

Units can be given biases by introducing an extra input to each

神经网络的历史

Brief History of Neural Networks

- The Third wave
 - 2006 **Deep (neural networks) Learning** gains popularity and
 - 2012 made significant break-through in many applications.
 - AlexNet—ImageNet champion model



Geoffery Hinton



Microsoft DNN research



神经网络的历史

Brief History of Neural Networks

- AlphaGo wins Lee Sedol (4-1)




<https://deepmind.com/research/alphago/>

Rank	Name	名字	Flag	Elo
1	Ke Jie	柯洁		3628
2	AlphaGo			3598
3	Park Junghwan	朴廷桓		3585
4	Tao Xiaxi	陶忻		3535
5	Mi Yuting	米昱婷		3534
6	Iyama Yuta	井山裕太		3525
7	Shi Yue	时越		3522
8	Lee Sedol	李昌镐		3521
9	Zhou Ruiyang	周睿羊		3517
10	Shin Jinseo	申真谞		3503
11	Chen Yaoye	陈耀烨		3495
12	Lian Xiao	连笑		3493
13	Tan Xiao	檀啸		3489
14	Kim Jiseok	金志锡		3489
15	Choi Cheolhan	崔哲瀚		3482
16	Park Yeonghwan	朴永桓		3482
17	Gu Zihao	古力		3468
18	Fan Yunruo	范蕴若		3468
19	Huang Yunsong	黄云嵩		3467
20	Li Qincheng	李钦诚		3465
21	Tang Weixing	唐韦星		3461
22	Lee Donghoon	李东勋		3460
23	Lee Yeongkyu	李映九		3459
24	Fan Tingyu	樊廷玉		3459
25	Tong Mengcheng	童梦成		3447
26	Kang Dongyun	康东润		3442
27	Yang Xi	杨鼎新		3439
28	Voon Seonjin	文相珍		3439
29	Yang Dingxin	杨丁欣		3439
30	Gu Li	古力		3436

<https://www.goratings.org/>


Breaking News of AI in 2016

大脑中的神经元 Neurons in the Brain



- The brain is composed of a **mass of interconnected neurons** (about 10^{11} neurons with an average of 10^4 connections each).
 - each neuron is connected to many other neurons



大脑中的神经元 Neurons in the Brain




- A neuron receives input from other neurons from its synapses
- Inputs are approximately summed
- When the input exceeds a threshold the neuron sends an electrical spike that travels from the body, down the axon, to the next neuron(s)

[Credit: US National Institutes of Health, National Institute on Aging]


人工神经元 Artificial Neuron Diagram



人工神经元 Artificial Neuron




$$\text{sgn}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$



感知器及其训练法则

Perceptron & Training Rule



Training

$$w_i = w_i + \eta(y - \hat{y}) x_i$$

$$b = b + \eta(y - \hat{y})$$

where η is the “learning rate”

Prediction

$$y = \sigma\left(\sum w_i x_i - b\right)$$

Activation function

$$\sigma(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$


- Equivalent to rules:
 - If output is correct do nothing.
 - If output is high, lower weights on active inputs
 - If output is low, increase weights on active inputs

感知器及其训练法则

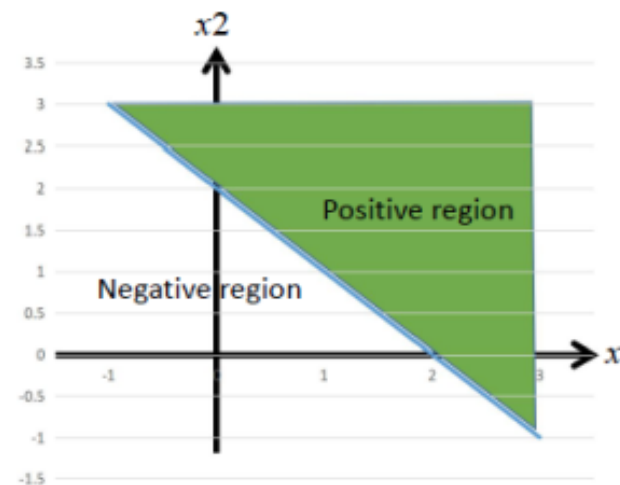
Perceptron & Training Rule

- Since perceptron uses linear threshold function, it is searching for a linear separator that discriminates the classes.

Perceptron



$$y = w_1x_1 + w_2x_2 + b$$




$$w_1 = 1, w_2 = 1, b = -2$$


Eric

感知器及其局限 Perceptron Limitation


- Can prove it will converge
 - If training data is linearly separable
 - And η sufficiently small
- However, a single layer perceptron can only learn linearly separable concepts.



(a) $AND (x_1 \cap x_2)$



(b) $OR (x_1 \cup x_2)$




(c) *Exclusive-OR*
 $(x_1 \oplus x_2)$

感知器及其训练法则


Perceptron & Training Rule

- Since perceptron uses linear threshold function, it is searching for a linear separator that discriminates the classes.

Perceptron



$$y = w_1x_1 + w_2x_2 + b$$




$$w_1 = 1, w_2 = 1, b = -2$$

Eric

加入隐藏层的神经网络


Network with One Hidden Layer

Perceptron



linear combination of three decision lines

single layer perceptron is a linear classifier



$$w_{1-11} = 1, w_{1-12} = 1, b_{1-1} = -2$$

$$w_{1-21} = 2, w_{1-22} = 1, b_{1-2} = 4$$


$$w_{1-31} = 0, w_{1-32} = 1, b_{1-3} = 3$$

https://blog.csdn.net/VinceVj_Eric

加入隐藏层的神经网络

Network with One Hidden Layer


- All Boolean functions can be represented with two-layers network.



加入隐藏层的神经网络


Network with One Hidden Layer

Perceptron



linear combination of three decision lines

single layer perceptron is a linear classifier




$$w_{1-11} = 1, w_{1-12} = 1, b_{1-1} = -2$$


$$w_{1-21} = 2, w_{1-22} = 1, b_{1-2} = 4$$

$$w_{1-31} = 0, w_{1-32} = 1, b_{1-3} = 3$$

人工神经元 Artificial Neuron



$$\text{sgn}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$




激活函数选择

Activation Function

- Activation function must be **continuous, differentiable, non-linear, and easy to compute**

Logistic Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Its derivative:

$$\sigma(z)' = ?$$


激活函数选择

Activation Function

- Activation function must be **continuous, differentiable, non-linear, and easy to compute**


Logistic Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Its derivative:

$$\sigma(z)' = \frac{e^{-z}}{(1 + e^{-z})^2}$$




激活函数选择

Activation Function

- Activation function must be **continuous, differentiable, non-linear, and easy to compute**


Logistic Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Its derivative:


$$\sigma(z)' = \frac{e^{-z}}{(1 + e^{-z})^2}$$




Output range [0,1];

Motivated by biological neurons and can be interpreted as the probability of an artificial neuron 'firing' given its inputs


人工神经元 Artificial Neuron




$$\text{sgn}(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$



人工神经元 Artificial Neuron




$$\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$




非线性激活函数 Non-linear activation function

Perceptron with non-linear activation f'




$$\begin{aligned} a1 &= w_{1-11}x_1 + w_{1-21}x_2 + b_{1-1} \\ a2 &= w_{1-12}x_1 + w_{1-22}x_2 + b_{1-2} \\ a3 &= w_{1-13}x_1 + w_{1-23}x_2 + b_{1-3} \end{aligned}$$

$$y = \sigma(w_{2-1}\sigma(a1) + w_{2-2}\sigma(a2) + w_{2-3}\sigma(a3))$$



with step activation function



with sigmoid activation function

神经网络的标准结构

Standard structure of an ANN

- **Input units**

- represents the input as a fixed-length vector of numbers (user defined)

- **Hidden units**

- represent intermediate calculations that the network learns;
- Usually just **one** (i.e., a 2-layer net)


- **Output units**

- represent the output as a fixed length vector of numbers

- *Units are connected by **links**.*

- *Each link has a **numeric weight**.*

多层神经网络 Multi-layers Neural Network



Network parameter θ : all the weights and biases in the "neurons"

Different connection leads to different network structures

- Given the training dataset of (data,label) pairs,

$$D = \{(x^{(i)}, y^{(i)})\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label

$$y^{(i)} \approx \boxed{f_{\theta}(x^{(i)})}$$

- Function set $\{f_{\theta}(x^{(i)})\}$ is called hypothesis space
- Learning is referred to as updating the parameter θ to make the prediction closed to the corresponding label

- Given the training dataset of (data,label) pairs,

$$D = \{(x^{(i)}, y^{(i)})\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label


$$y^{(i)} \approx f_{\theta}(x^{(i)}) \rightarrow \text{neural network}$$

- Function set $\{f_{\theta}(x^{(i)})\}$ is called hypothesis set
- Learning is referred to as updating the parameters θ so that the prediction is as close to the corresponding label




prediction


全连接前馈神经网络 Fully Connect Feedforward Network




全连接前馈神经网络 Fully Connect Feedforward Network



全连接前馈神经网络 Fully Connect Feedforward Network



全连接前馈神经网络 Fully Connect Feedforward Network




This is a function.

Input vector, output vector

$$f\left(\begin{bmatrix} 1 \\ -1 \end{bmatrix}\right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix} \quad f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$$

Given network structure, define a function set

全连接前馈神经网络 Fully Connect Feedforward Network



This is a function.

Input vector, output vector

$$f\left(\begin{bmatrix} 1 \\ -1 \end{bmatrix}\right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix} \quad f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$$

Given network structure, define a function set

Different w and bias

问题1：学习目标 What is the Learning Objective?

- Make the prediction closed to the corresponding label

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N L(y^{(i)}, f_{\theta}(x^{(i)}))$$

(Empirical Risk Minimization, ERM)

Loss function $L(y^{(i)}, f_{\theta}(x^{(i)}))$ measures the error between the label and prediction for single sample.

The definition of loss function depends on the data and task

损失函数 Loss function

0-1 Loss Function

$$L(y^{(i)}, f(x^{(i)})) = \begin{cases} 1, & \text{if } y^{(i)} \neq f(x^{(i)}) \\ 0, & \text{if } y^{(i)} = f(x^{(i)}) \end{cases}$$

Mean Squared Error, MSE

$$L(y^{(i)}, f(x^{(i)})) = (y^{(i)} - f(x^{(i)}))^2$$

Absolute Loss Function


$$L(y^{(i)}, f(x^{(i)})) = |y^{(i)} - f(x^{(i)})|$$

Logarithmic Loss Function (Cross-Entropy Loss Function)

$$L(y^{(i)}, p^{(i)}) = -[y^{(i)} \log(p^{(i)}) + (1 - y^{(i)}) \log(1 - p^{(i)})]$$

$y^{(i)} \in \{0, 1\}$, $p^{(i)} = f(x^{(i)})$ is the predicted probability that the i th sample belongs to the positive class (usually denoted as class 1))

全连接前馈神经网络 Fully Connect Feedforward Network



梯度下降求解

Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$

Starting Parameters $\theta^0 \longrightarrow \theta^1 \longrightarrow \theta^2 \longrightarrow \dots$


$$\nabla L(\theta) = \begin{bmatrix} \partial L(\theta) / \partial w_1 \\ \partial L(\theta) / \partial w_2 \\ \vdots \\ \partial L(\theta) / \partial b_1 \\ \partial L(\theta) / \partial b_2 \\ \vdots \end{bmatrix}$$

$\text{Compute } \nabla L(\theta^0) \quad \theta^1 = \theta^0 - \eta \nabla L(\theta^0)$
 $\text{Compute } \nabla L(\theta^1) \quad \theta^2 = \theta^1 - \eta \nabla L(\theta^1)$
 Millions of parameters
 To compute the gradients efficiently,
 we use backpropagation.


反向传播算法 Backpropagation

Cost function: $\frac{1}{N} \sum_{i=1}^N L(y^{(i)}, y'^{(i)})$

Regression:



Classification:



Class 1 Class 2 Class 3


$\hat{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$ or $\begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$ or $\begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$

Class as One hot vector

反向传播算法 Backpropagation


Cost function: $\frac{1}{N} \sum_{i=1}^N L(y^{(i)}, y'^{(i)})$

Regression:



Mean Square Error (MSE): $L(y', y) = (y' - y)^2$

Classification:



$L(y', y) = (y' - y)^2$

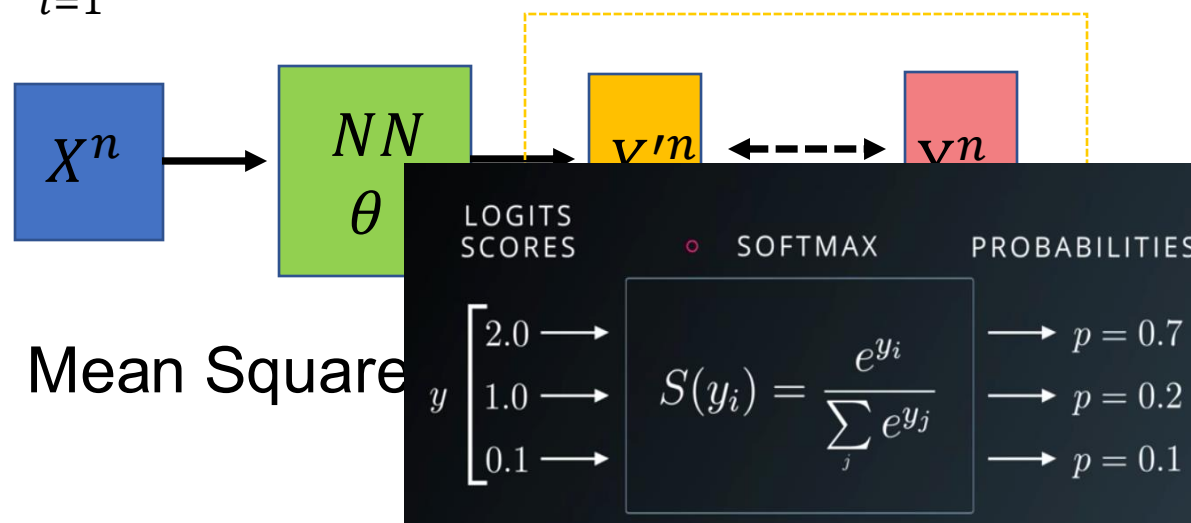
Class as One hot vector

Cross-entropy: $L(y', y) = - \sum y \ln y'$

反向传播算法 Backpropagation


Cost function: $\frac{1}{N} \sum_{i=1}^N L(y^{(i)}, y'^{(i)})$

Regression:



Mean Square

Classification:



$$L(y', y) = (y' - y)^2$$

Class as One hot vector

Cross-entropy: $L(y', y) = - \sum y \ln y'$

Softmax回归 Softmax Regression

- Input: $x = (x_1, x_2, \dots, x_n)^T$
- Output: class label $\in \{0, 1, \dots, K\}$ represented by one-hot vector

$$y = (I(1 = k), I(2 = k), \dots, I(K = k))^T$$

- For each class k , there is a weight vector θ_k
- For each class k , the predicted probability is : $p = \frac{e^{\theta_k^T x}}{\sum_{j=1}^K e^{\theta_j^T x}}$
- Using cross-entropy, the cost function:

$$-\sum_{i=1}^N \sum_{j=1}^K y_j^i \log(p_j^i)$$

- Predicting label:

$$\hat{y} = \arg \max_{j=1 \dots K} p_j(x)$$

Softmax回归 Softmax Regression

- Input: $x = (x_1, x_2, \dots, x_n)^T$
- Output: class label $\in \{0, 1, \dots, K\}$ represented by one-hot vector

$$y = (I(1 = k), I(2 = k), \dots, I(K = k))^T$$

e.g. $K=3$, then $label1 = (1, 0, 0)^T$
 $label2 = (0, 1, 0)^T$
 $label3 = (0, 0, 1)^T$

- For each class k , there is a weight vector θ_k

- For each class k , the predicted probability is : $p = \frac{e^{\theta_k^T x}}{\sum_{j=1}^K e^{\theta_j^T x}}$

- Using cross-entropy, the cost function:

$$-\sum_{i=1}^N \sum_{j=1}^K y_j^i \log(p_j^i)$$

e.g. $y = (0, 0, 1)^T, \quad p = (0.3, 0.3, 0.4)^T$
 $loss = -(0 * \log(0.3) + 0 * \log(0.3) + 1 * \log(0.4))$
 ≈ 0.916

- Predicting label:


$$\hat{y} = \arg \max_{j=1 \dots K} p_j(x)$$

$y = (0, 0, 1)^T, \quad p = (0.1, 0.1, 0.8)^T$
 $loss = -(0 * \log(0.1) + 0 * \log(0.1) + 1 * \log(0.8))$
 ≈ 0.223

反向传播算法 Backpropagation


Cost function: $\frac{1}{N} \sum_{i=1}^N L(y^{(i)}, y'^{(i)})$

Regression:



Mean Square Error (MSE): $L(y', y) = (y' - y)^2$

Classification:




$$L(y', y) = (y' - y)^2$$

Class as One hot vector

Cross-entropy: $L(y', y) = - \sum y \ln y'$

反向传播算法 Backpropagation



Forward pass:

Compute $\partial z / \partial w$ for all parameters

Backward pass:


Compute $\partial C / \partial z$ for all activation function inputs z

$$\frac{\partial C}{\partial w} = ? \quad \frac{\partial z}{\partial w} \frac{\partial C}{\partial z}$$

(Chain rule)

反向传播算法之正向传递 Backpropagation-Forward Pass

Compute $\partial z / \partial w$ for all parameters




$$\left. \begin{aligned} \partial z / \partial w_1 &= ? \quad x_1 \\ \partial z / \partial w_2 &= ? \quad x_2 \end{aligned} \right\}$$

The value of the input
connected by the weight

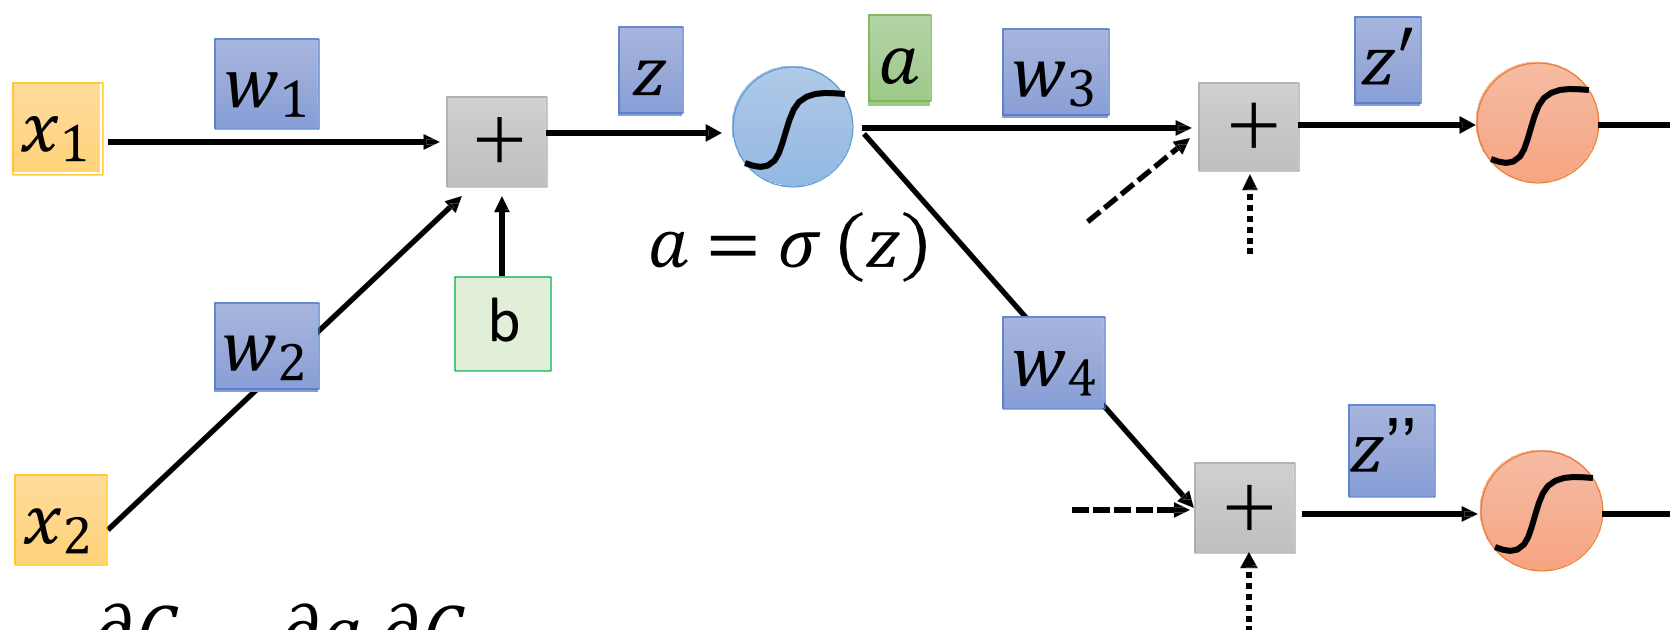
反向传播算法之正向传递 Backpropagation-Forward Pass

Compute $\partial z / \partial w$ for all parameters



反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z

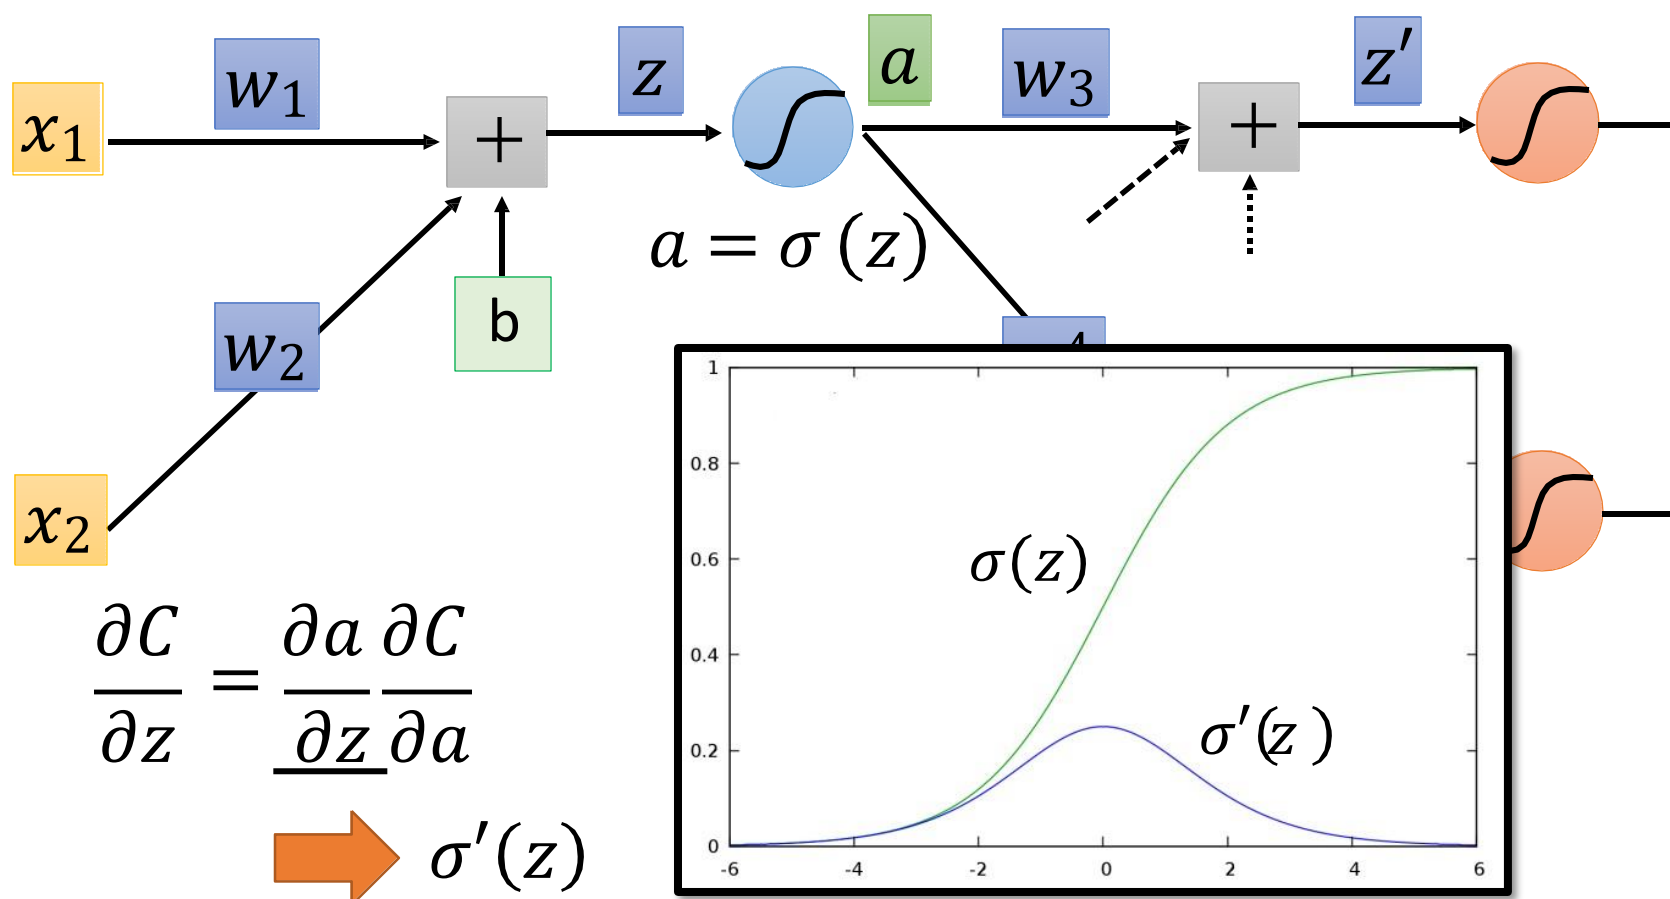


$$\frac{\partial C}{\partial z} = \frac{\partial a}{\partial z} \frac{\partial C}{\partial a}$$

➡ $\sigma'(z)$

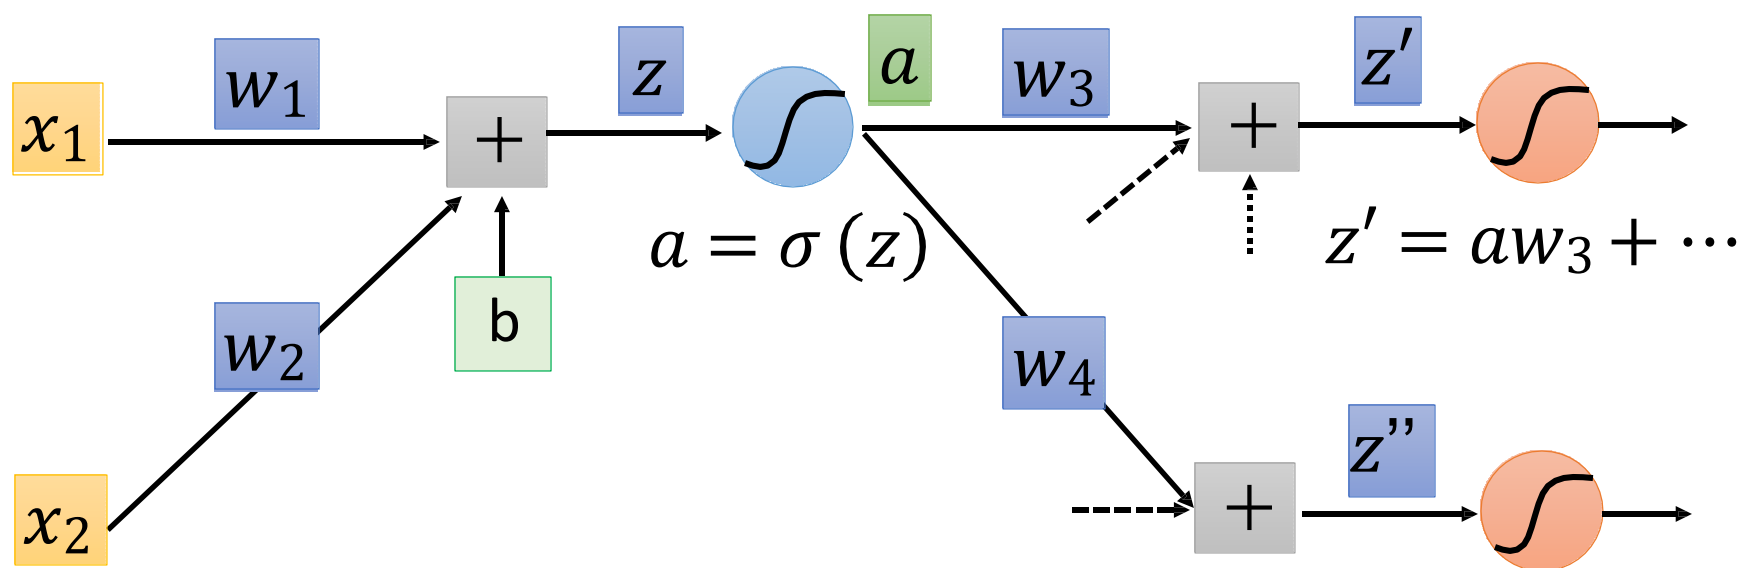
反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z



反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z



$$\frac{\partial C}{\partial z} = \frac{\partial a}{\partial z} \frac{\partial C}{\partial a} \quad \frac{\partial C}{\partial a} = \frac{\partial z'}{\partial a} \frac{\partial C}{\partial z'} + \frac{\partial z''}{\partial a} \frac{\partial C}{\partial z''} \quad (\text{Chain rule})$$

w_3

?

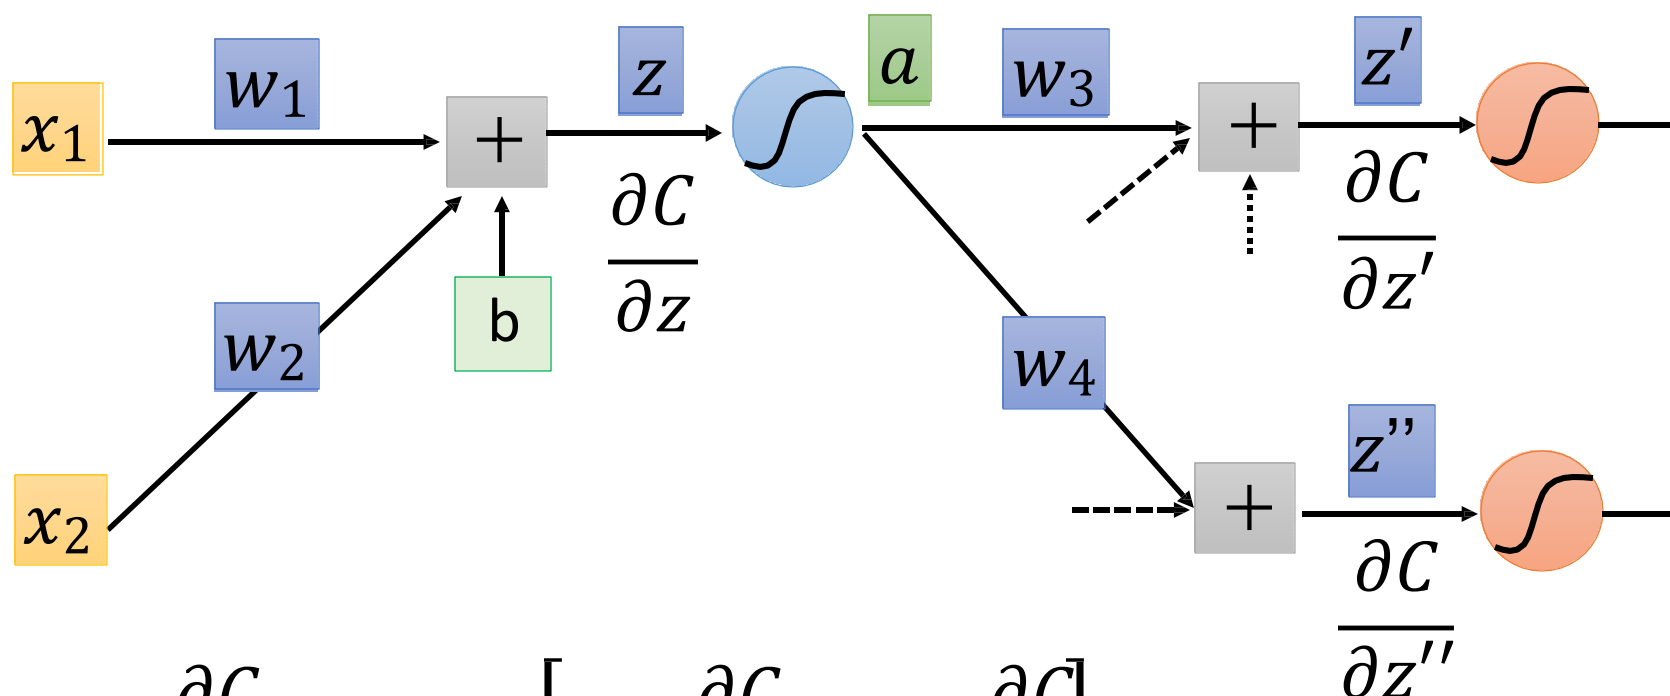
w_4

?

Assumed
it's known

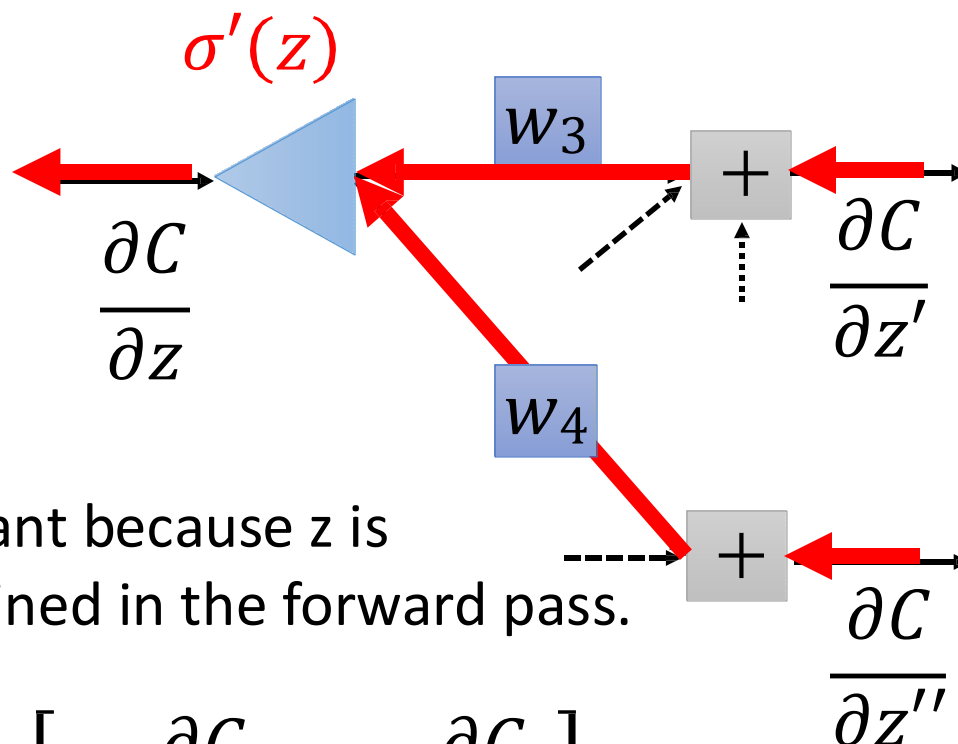
反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z



$$\frac{\partial C}{\partial z} = \sigma'(z) \left[w_3 \frac{\partial C}{\partial z'} + w_4 \frac{\partial C}{\partial z''} \right]$$

反向传播算法-反向传递 Backpropagation-backward Pass

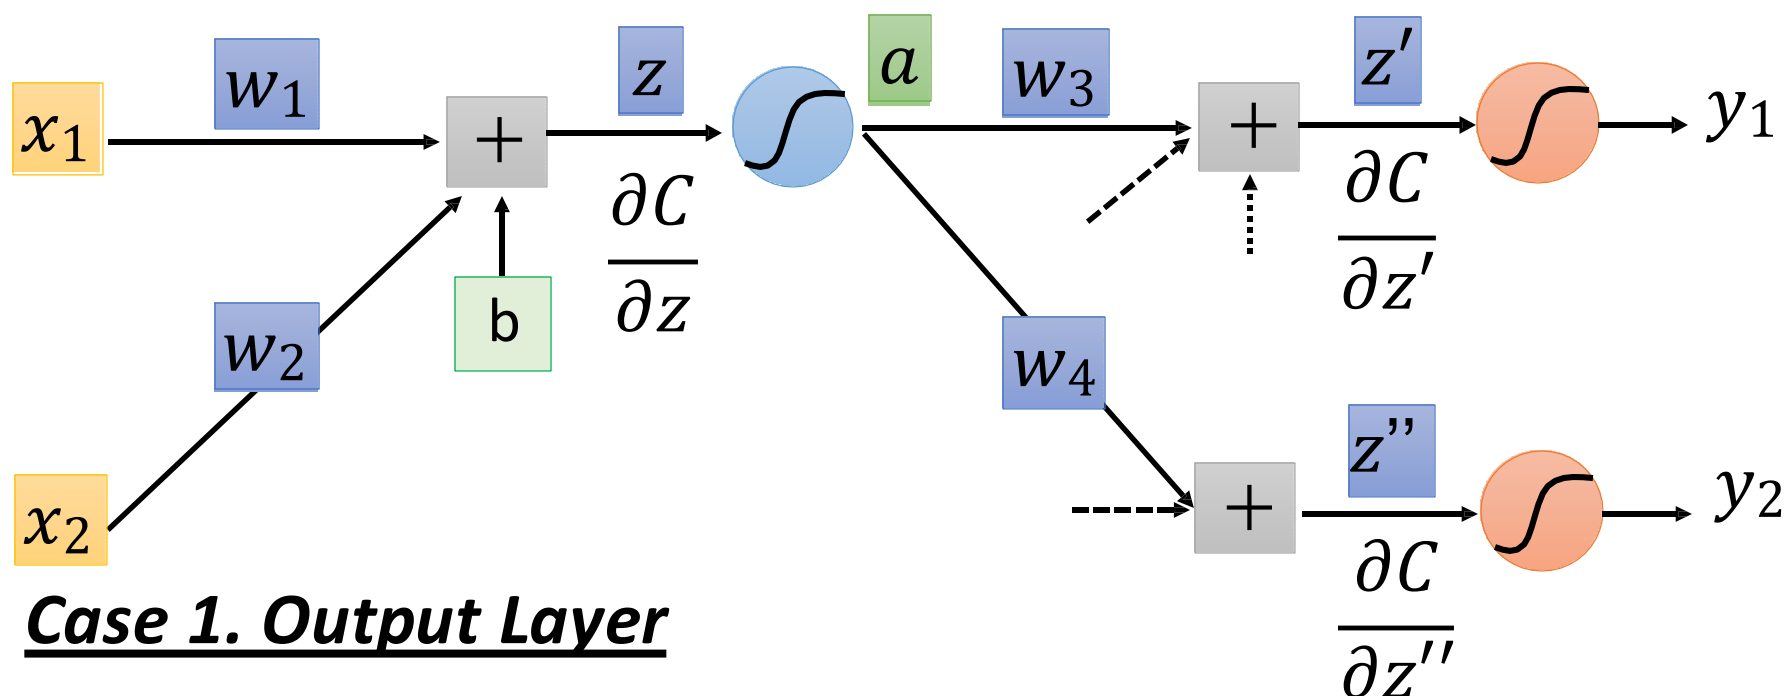


$\sigma'(z)$ is a constant because z is already determined in the forward pass.

$$\frac{\partial C}{\partial z} = \sigma'(z) \left[w_3 \frac{\partial C}{\partial z'} + w_4 \frac{\partial C}{\partial z''} \right]$$

反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z




$$\frac{\partial C}{\partial z'} = \frac{\partial y_1}{\partial z'} \frac{\partial C}{\partial y_1} \quad \frac{\partial C}{\partial z''} = \frac{\partial y_2}{\partial z''} \frac{\partial C}{\partial y_2}$$

Done!

反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z

Case 2. Not Output Layer

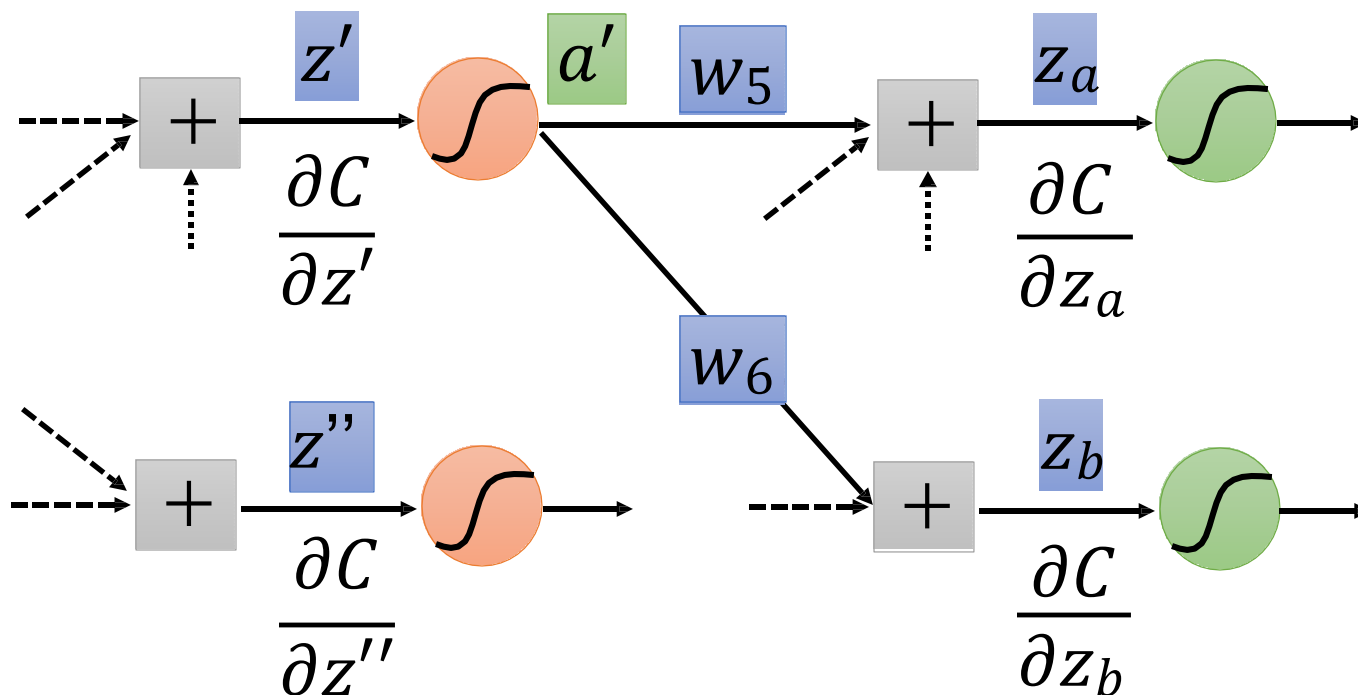


反向传播算法-反向传递

Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z

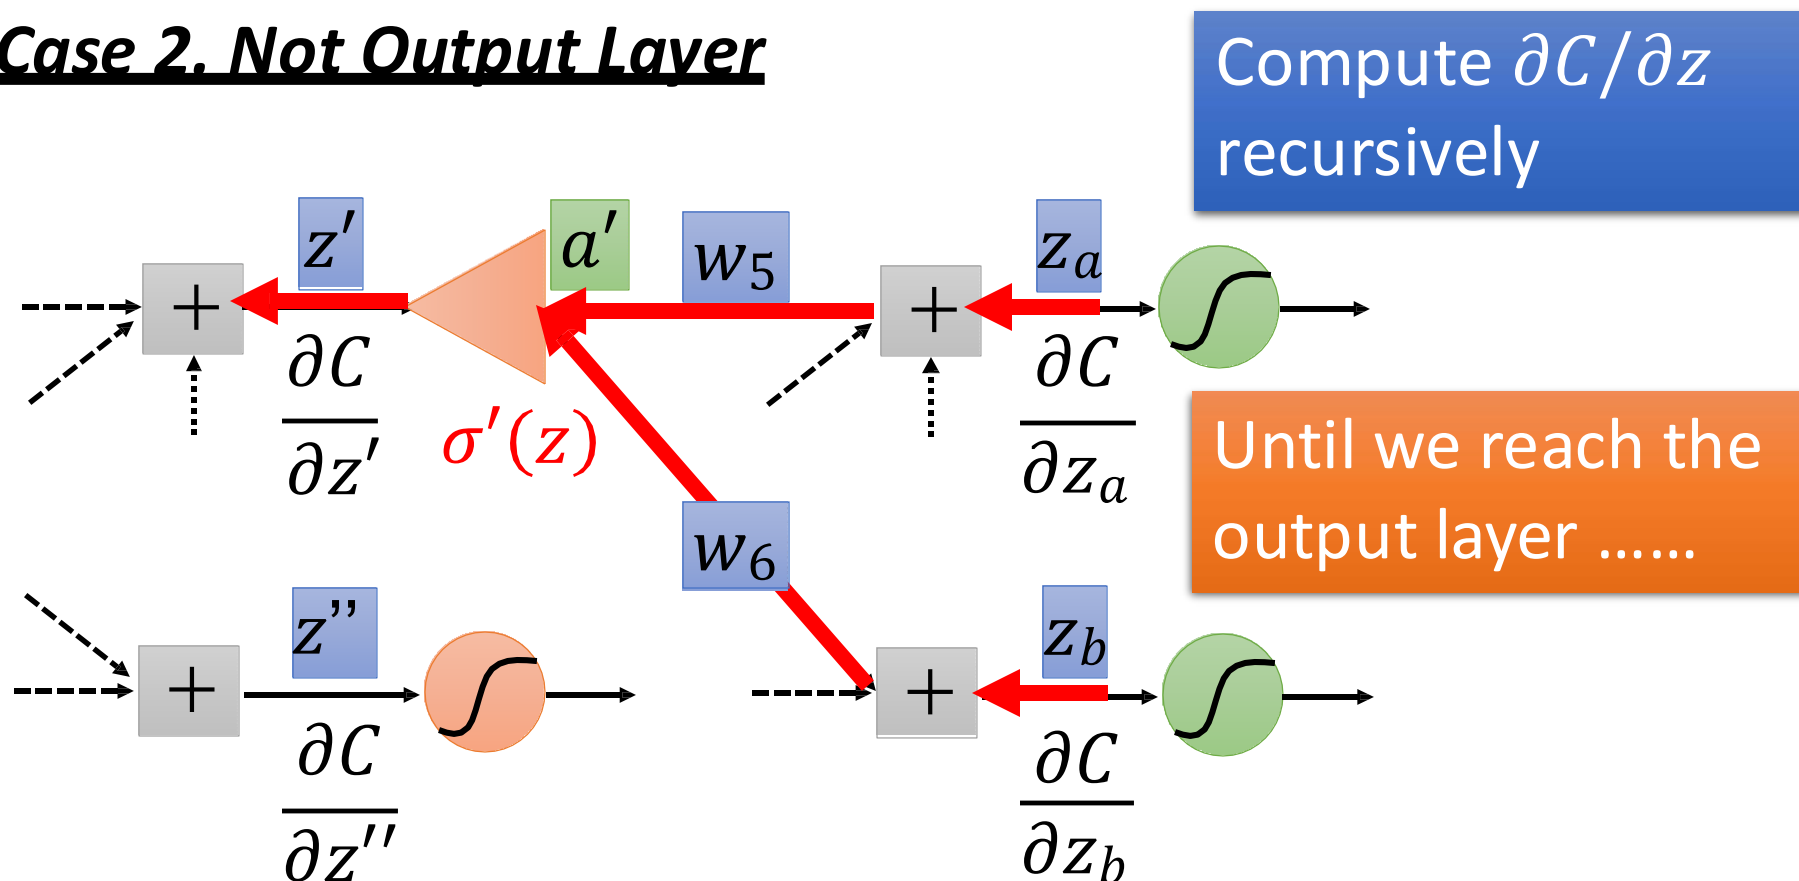
Case 2. Not Output Layer



反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z


Case 2. Not Output Layer



反向传播算法-反向传递 Backpropagation-backward Pass

Compute $\partial C / \partial z$ for all activation function inputs z

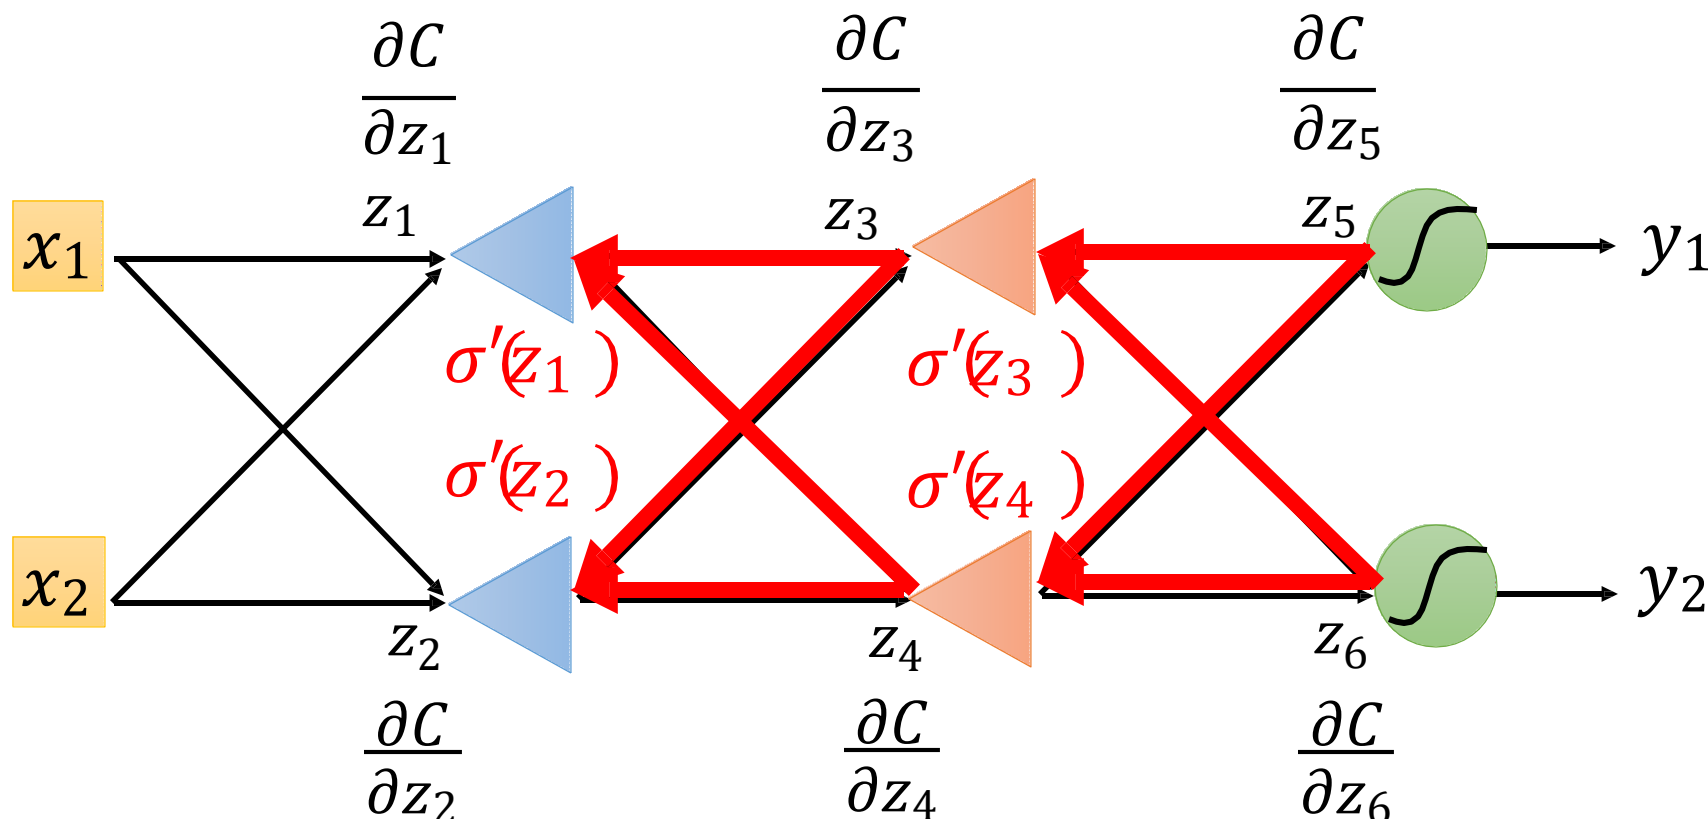
Compute $\partial C / \partial z$ from the output layer



反向传播算法-反向传递 Backpropagation-backward Pass


Compute $\partial C / \partial z$ for all activation function inputs z

Compute $\partial C / \partial z$ from the output layer




反向传播法 Backpropagation

Forward Pass



$$\frac{\partial z}{\partial w} = a$$

Backward Pass



X

$$\frac{\partial C}{\partial z}$$


$$= \frac{\partial C}{\partial w}$$

for all w

反向传播法 Backpropagation


Cost function: $\frac{1}{N} \sum_{i=1}^N L(y^{(i)}, y'^{(i)})$

Regression:



Mean Square Error (MSE): $L(y', y) = (y' - y)^2$


Classification:



CSDN @特快雪时晴


反向传播法 Backpropagation

Regression:



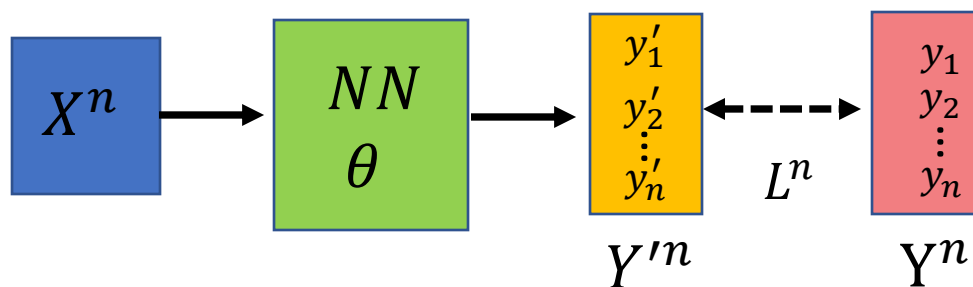
Mean Square Error (MSE): $L(y', y) = (y' - y)^2$

Classification:



Cross-entropy: $L(y', y) = - \sum y \ln y'$

反向传播法 Backpropagation




MSE:
$$L(Y', Y) = \frac{1}{2} \sum_i^N (y'_i - y_i)^2$$

$$\frac{\partial L}{\partial y'_i} == y'_i - y_i$$


反向传播法 Backpropagation

Regression:



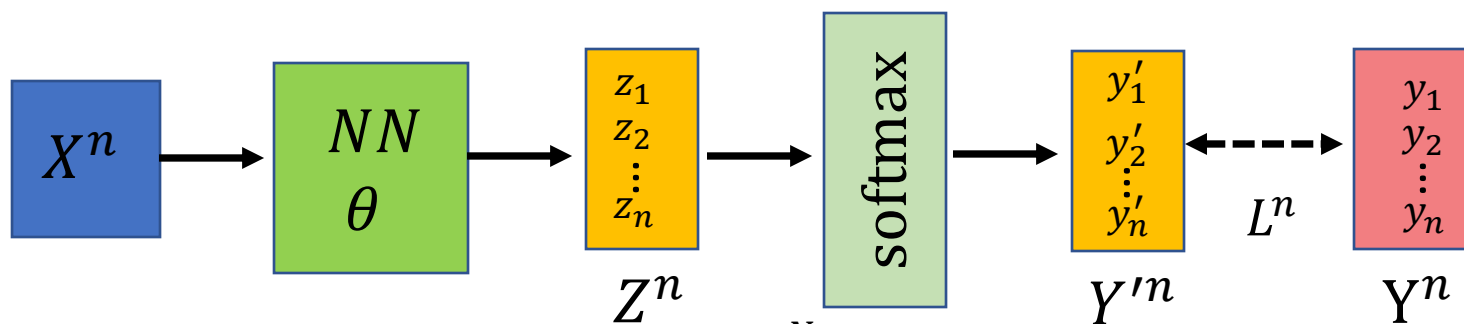
Mean Square Error (MSE): $L(y', y) = (y' - y)^2$

Classification:



Cross-entropy: $L(y', y) = - \sum y \ln y'$

反向传播法 Backpropagation



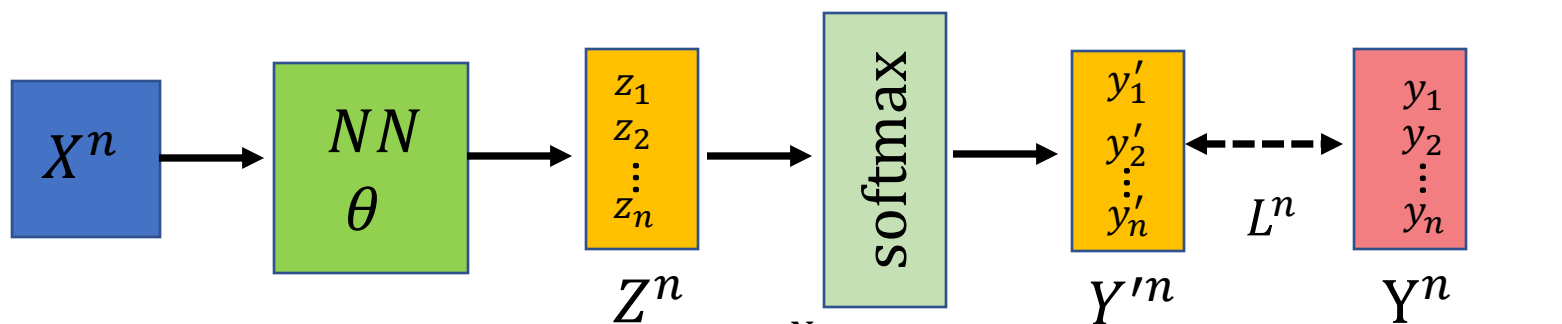
Cross-entropy: $L(Y', Y) = - \sum_i^N y_i \cdot \log y'_i$

$$y'_i = \frac{e^{z_i}}{\sum_{k=1}^N e^{z_k}}$$

$$\frac{\partial y'_j}{\partial z_i} = \frac{\partial \frac{e^{z_j}}{\sum_{k=1}^N e^{z_k}}}{\partial z_i} = \frac{\frac{\partial e^{z_j}}{\partial z_i} \cdot \sum_{k=1}^N e^{z_k} - \frac{\partial \sum_{k=1}^N e^{z_k}}{\partial z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2}$$

$$\begin{cases} \frac{-e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = -y'_i \cdot y'_j, & i \neq j \\ \frac{e^{z_i} \cdot \sum_{k=1}^N e^{z_k} - e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = y'_i - y'_i \cdot y'_j, & i = j \end{cases}$$

反向传播法 Backpropagation



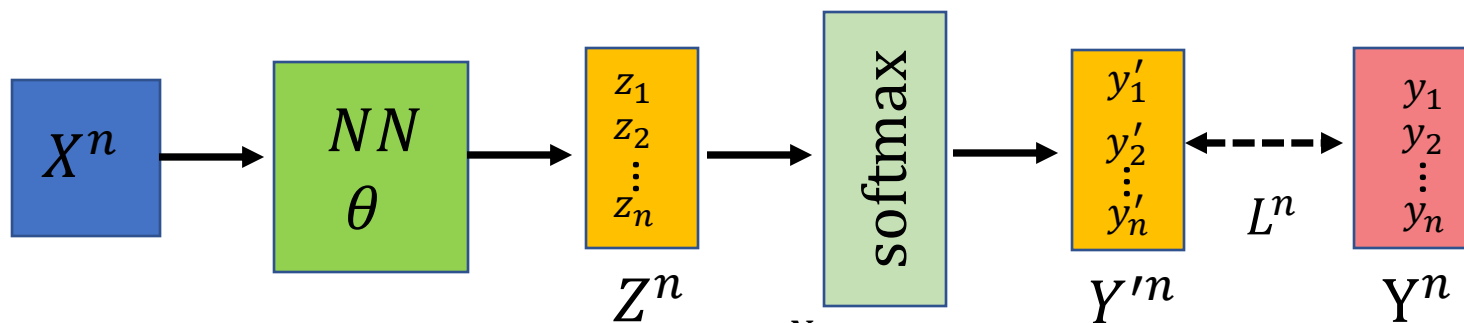
Cross-entropy: $L(Y', Y) = - \sum_i^N y_i \cdot \log y'_i$

$$y'_i = \frac{e^{z_i}}{\sum_{k=1}^N e^{z_k}}$$

$$\frac{\partial y'_j}{\partial z_i} =$$

$$\begin{cases} \frac{-e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = -y'_i \cdot y'_j, & i \neq j \\ \frac{e^{z_i} \cdot \sum_{k=1}^N e^{z_k} - e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = y'_i - y'_i \cdot y'_j, & i = j \end{cases}$$

反向传播法 Backpropagation




Cross-entropy: $L(Y', Y) = - \sum_i^N y_i \cdot \log y'_i$

$$y'_i = \frac{e^{z_i}}{\sum_{k=1}^N e^{z_k}}$$

$$\frac{\partial y'_j}{\partial z_i} = \begin{cases} \frac{-e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = -y'_i \cdot y'_j, & i \neq j \\ \frac{e^{z_i} \cdot \sum_{k=1}^N e^{z_k} - e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = y'_i - y'_i \cdot y'_j, & i = j \end{cases}$$

$$\frac{\partial L}{\partial z_i} = \sum_{j=1}^N \frac{\partial L}{\partial y'_j} \frac{\partial y'_j}{\partial z_i}$$

反向传播法 Backpropagation



Cross-entropy: $L(Y', Y) = - \sum_i y_i \cdot \log y'_i$ $y'_i = \frac{e^{z_i}}{\sum_{k=1}^N e^{z_k}}$

$$\frac{\partial y'_j}{\partial z_i} = \begin{cases} \frac{-e^{z_i} \cdot e^{z_j}}{(\sum_{k=1}^N e^{z_k})^2} = -y'_i \cdot y'_j, & i \neq j \\ \frac{e^{z_i} \cdot \sum_{k=1, k \neq i}^N e^{z_k} - e^{z_i} \cdot e^{z_i}}{(\sum_{k=1}^N e^{z_k})^2} = y'_i - y'_i \cdot y'_i, & i = j \end{cases}$$

$$\frac{\partial L}{\partial z_i} = \sum_{j=1}^N \frac{\partial L}{\partial y'_j} \frac{\partial y'_j}{\partial z_i}$$


$$= -\frac{y_i}{y'_i} (y'_i - y'_i \cdot y'_i) - \sum_{j=1, j \neq i}^N \frac{y_j}{y'_j} (-y'_i \cdot y'_j)$$

$$= -y_i + y_i \cdot y'_i + \sum_{j=1, j \neq i}^N y_j \cdot y'_i$$

$$= -y_i + y'_i$$

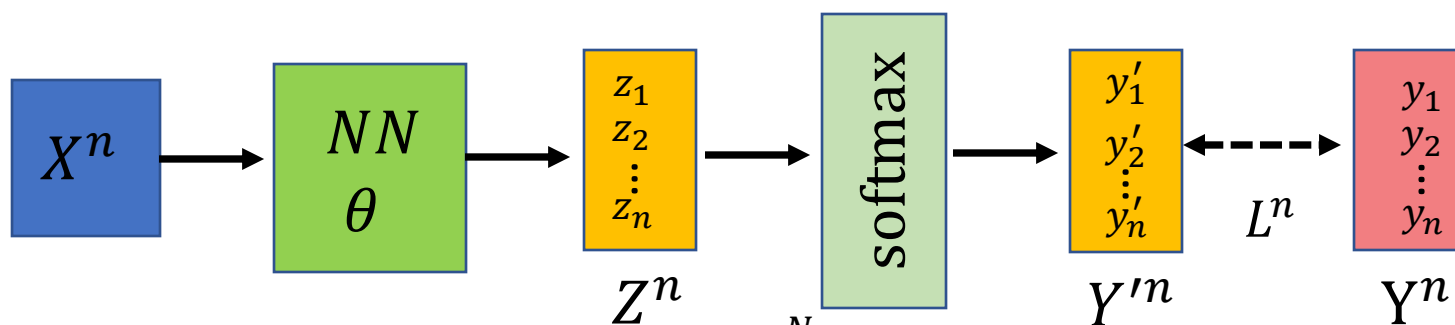
$$= y'_i - y_i$$

反向传播法 Backpropagation



MSE:
$$L(Y', Y) = \frac{1}{2} \sum_i^N (y'_i - y_i)^2$$

$$\frac{\partial L}{\partial y'_i} == y'_i - y_i$$



Cross-entropy:
$$L(Y', Y) = - \sum_i^N y_i \cdot \log y'_i$$

$$\frac{\partial L}{\partial z_i} == y'_i - y_i$$

反向传播算法评价

Remarks on the Backprop

- Very powerful - can learn any function with enough hidden units, we can generate any function.
- BP is only guaranteed to converge toward some local minimum and not necessarily to the global minimum error.
 - momentum term
 - stochastic gradient descent
 - Train multiple networks using the same data

多层神经网络的表征能力

Representational Power of ANN



同济大学
TONGJI UNIVERSITY

- **Boolean functions**
 - Any boolean function can be represented by a two-layer network with sufficient hidden units.
- **Continuous functions**
 - Any bounded continuous function can be approximated with arbitrarily small error by a two-layer network.
- **Arbitrary function**
 - Any function can be approximated to arbitrary accuracy by a three-layer network.

神经网络的建立 Establishment of Neural Networks

- Decisions must be taken on the following:
 - The number of units to use.(hyperparameters)
 - How many hidden units in the layer?
 - Too few ==> can't learn
 - Too many ==> poor generalization
 - The type of units required.(hyperparameters)
 - Connection between the units.

隐藏单元的特征

Hidden Unit Representations


- Trained hidden units can be seen as **newly constructed features** that make the target concept linearly separable in the transformed space.
- On many real domains, hidden units can be interpreted as **representing meaningful features** such as vowel detectors or edge detectors, etc..
- However, the hidden layer can also become a distributed representation of the input in which each individual unit is not easily interpretable as a meaningful feature.

隐藏单元的代表征举例

Example - Hidden Unit Representations

A target function:

Input	Output
10000000 →	10000000
01000000 →	01000000
00100000 →	00100000
00010000 →	00010000
00001000 →	00001000
00000100 →	00000100
00000010 →	00000010
00000001 →	00000001




Can this be learned by a 8*3*8 neural network?

<http://www.cs.cmu.edu/~tom/mlbook>

隐藏单元的特征举例

Example - Hidden Unit Representations



Input		Hidden Values				Output
10000000	→	.89	.04	.08	→	10000000
01000000	→	.15	.99	.99	→	01000000
00100000	→	.01	.97	.27	→	00100000
00010000	→	.99	.97	.71	→	00010000
00001000	→	.03	.05	.02	→	00001000
00000100	→	.01	.11	.88	→	00000100
00000010	→	.80	.01	.98	→	00000010
00000001	→	.60	.94	.01	→	00000001

Learned hidden layer representation after 5000 times training

<http://www.cs.cmu.edu/~tom/mlbook>

小结 Recap

- Neural network is a computational model that **simulate** some properties of the **human brain**.
- The **connections and nature of units** determine the behavior of a neural network.
- Perceptrons are feed-forward networks that can only represent **linearly separable functions**.
- Given enough units, **any function can be represented** by Multi-layer feed-forward networks.
- **Backpropagation** learning works on multi-layer feed-forward networks.
- Neural Networks are widely used in developing artificial learning systems.